

# boston\_housing

May 22, 2019

## 1 Machine Learning Engineer Nanodegree

### 1.1 Model Evaluation & Validation

### 1.2 Project: Predicting Boston Housing Prices

Welcome to the first project of the Machine Learning Engineer Nanodegree! In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**Implementation**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

### 1.3 Getting Started

In this project, you will evaluate the performance and predictive power of a model that has been trained and tested on data collected from homes in suburbs of Boston, Massachusetts. A model trained on this data that is seen as a *good fit* could then be used to make certain predictions about a home — in particular, its monetary value. This model would prove to be invaluable for someone like a real estate agent who could make use of such information on a daily basis.

The dataset for this project originates from the [UCI Machine Learning Repository](#). The Boston housing data was collected in 1978 and each of the 506 entries represent aggregated data about 14 features for homes from various suburbs in Boston, Massachusetts. For the purposes of this project, the following preprocessing steps have been made to the dataset: - 16 data points have an 'MEDV' value of 50.0. These data points likely contain **missing or censored values** and have been removed. - 1 data point has an 'RM' value of 8.78. This data point can be considered an

**outlier** and has been removed. - The features 'RM', 'LSTAT', 'PTRATIO', and 'MEDV' are essential. The remaining **non-relevant features** have been excluded. - The feature 'MEDV' has been **multiplicatively scaled** to account for 35 years of market inflation.

Run the code cell below to load the Boston housing dataset, along with a few of the necessary Python libraries required for this project. You will know the dataset loaded successfully if the size of the dataset is reported.

```
In [2]: import warnings
        warnings.filterwarnings("ignore", category=DeprecationWarning)

        # Import libraries necessary for this project
        import numpy as np
        import pandas as pd
        from sklearn.cross_validation import ShuffleSplit

        # Import supplementary visualizations code visuals.py
        import visuals as vs

        # Pretty display for notebooks
        %matplotlib inline

        # Load the Boston housing dataset
        data = pd.read_csv('housing.csv')
        prices = data['MEDV']
        features = data.drop('MEDV', axis = 1)

        # Success
        print("Boston housing dataset has {} data points with {} variables each.".format(*data.s
```

Boston housing dataset has 489 data points with 4 variables each.

```
In [3]: data
```

```
Out[3]:
```

	RM	LSTAT	PTRATIO	MEDV
0	6.575	4.98	15.3	504000.0
1	6.421	9.14	17.8	453600.0
2	7.185	4.03	17.8	728700.0
3	6.998	2.94	18.7	701400.0
4	7.147	5.33	18.7	760200.0
5	6.430	5.21	18.7	602700.0
6	6.012	12.43	15.2	480900.0
7	6.172	19.15	15.2	569100.0
8	5.631	29.93	15.2	346500.0
9	6.004	17.10	15.2	396900.0
10	6.377	20.45	15.2	315000.0
11	6.009	13.27	15.2	396900.0
12	5.889	15.71	15.2	455700.0
13	5.949	8.26	21.0	428400.0

14	6.096	10.26	21.0	382200.0
15	5.834	8.47	21.0	417900.0
16	5.935	6.58	21.0	485100.0
17	5.990	14.67	21.0	367500.0
18	5.456	11.69	21.0	424200.0
19	5.727	11.28	21.0	382200.0
20	5.570	21.02	21.0	285600.0
21	5.965	13.83	21.0	411600.0
22	6.142	18.72	21.0	319200.0
23	5.813	19.88	21.0	304500.0
24	5.924	16.30	21.0	327600.0
25	5.599	16.51	21.0	291900.0
26	5.813	14.81	21.0	348600.0
27	6.047	17.28	21.0	310800.0
28	6.495	12.80	21.0	386400.0
29	6.674	11.98	21.0	441000.0
..	...	...	...	...
459	6.484	18.68	20.2	350700.0
460	5.304	24.91	20.2	252000.0
461	6.185	18.03	20.2	306600.0
462	6.229	13.11	20.2	449400.0
463	6.242	10.74	20.2	483000.0
464	6.750	7.74	20.2	497700.0
465	7.061	7.01	20.2	525000.0
466	5.762	10.42	20.2	457800.0
467	5.871	13.34	20.2	432600.0
468	6.312	10.58	20.2	445200.0
469	6.114	14.98	20.2	401100.0
470	5.905	11.45	20.2	432600.0
471	5.454	18.06	20.1	319200.0
472	5.414	23.97	20.1	147000.0
473	5.093	29.68	20.1	170100.0
474	5.983	18.07	20.1	285600.0
475	5.983	13.35	20.1	422100.0
476	5.707	12.01	19.2	457800.0
477	5.926	13.59	19.2	514500.0
478	5.670	17.60	19.2	485100.0
479	5.390	21.14	19.2	413700.0
480	5.794	14.10	19.2	384300.0
481	6.019	12.92	19.2	445200.0
482	5.569	15.10	19.2	367500.0
483	6.027	14.33	19.2	352800.0
484	6.593	9.67	21.0	470400.0
485	6.120	9.08	21.0	432600.0
486	6.976	5.64	21.0	501900.0
487	6.794	6.48	21.0	462000.0
488	6.030	7.88	21.0	249900.0

[489 rows x 4 columns]

In [4]: prices

```
Out[4]: 0      504000.0
        1      453600.0
        2      728700.0
        3      701400.0
        4      760200.0
        5      602700.0
        6      480900.0
        7      569100.0
        8      346500.0
        9      396900.0
       10      315000.0
       11      396900.0
       12      455700.0
       13      428400.0
       14      382200.0
       15      417900.0
       16      485100.0
       17      367500.0
       18      424200.0
       19      382200.0
       20      285600.0
       21      411600.0
       22      319200.0
       23      304500.0
       24      327600.0
       25      291900.0
       26      348600.0
       27      310800.0
       28      386400.0
       29      441000.0
       ...
      459      350700.0
      460      252000.0
      461      306600.0
      462      449400.0
      463      483000.0
      464      497700.0
      465      525000.0
      466      457800.0
      467      432600.0
      468      445200.0
      469      401100.0
      470      432600.0
      471      319200.0
```

```
472    147000.0
473    170100.0
474    285600.0
475    422100.0
476    457800.0
477    514500.0
478    485100.0
479    413700.0
480    384300.0
481    445200.0
482    367500.0
483    352800.0
484    470400.0
485    432600.0
486    501900.0
487    462000.0
488    249900.0
Name: MEDV, Length: 489, dtype: float64
```

```
In [5]: features
```

```
Out[5]:
```

	RM	LSTAT	PTRATIO
0	6.575	4.98	15.3
1	6.421	9.14	17.8
2	7.185	4.03	17.8
3	6.998	2.94	18.7
4	7.147	5.33	18.7
5	6.430	5.21	18.7
6	6.012	12.43	15.2
7	6.172	19.15	15.2
8	5.631	29.93	15.2
9	6.004	17.10	15.2
10	6.377	20.45	15.2
11	6.009	13.27	15.2
12	5.889	15.71	15.2
13	5.949	8.26	21.0
14	6.096	10.26	21.0
15	5.834	8.47	21.0
16	5.935	6.58	21.0
17	5.990	14.67	21.0
18	5.456	11.69	21.0
19	5.727	11.28	21.0
20	5.570	21.02	21.0
21	5.965	13.83	21.0
22	6.142	18.72	21.0
23	5.813	19.88	21.0
24	5.924	16.30	21.0
25	5.599	16.51	21.0

26	5.813	14.81	21.0
27	6.047	17.28	21.0
28	6.495	12.80	21.0
29	6.674	11.98	21.0
...	...	...	...
459	6.484	18.68	20.2
460	5.304	24.91	20.2
461	6.185	18.03	20.2
462	6.229	13.11	20.2
463	6.242	10.74	20.2
464	6.750	7.74	20.2
465	7.061	7.01	20.2
466	5.762	10.42	20.2
467	5.871	13.34	20.2
468	6.312	10.58	20.2
469	6.114	14.98	20.2
470	5.905	11.45	20.2
471	5.454	18.06	20.1
472	5.414	23.97	20.1
473	5.093	29.68	20.1
474	5.983	18.07	20.1
475	5.983	13.35	20.1
476	5.707	12.01	19.2
477	5.926	13.59	19.2
478	5.670	17.60	19.2
479	5.390	21.14	19.2
480	5.794	14.10	19.2
481	6.019	12.92	19.2
482	5.569	15.10	19.2
483	6.027	14.33	19.2
484	6.593	9.67	21.0
485	6.120	9.08	21.0
486	6.976	5.64	21.0
487	6.794	6.48	21.0
488	6.030	7.88	21.0

[489 rows x 3 columns]

## 1.4 Data Exploration

In this first section of this project, you will make a cursory investigation about the Boston housing data and provide your observations. Familiarizing yourself with the data through an explorative process is a fundamental practice to help you better understand and justify your results.

Since the main goal of this project is to construct a working model which has the capability of predicting the value of houses, we will need to separate the dataset into **features** and the **target variable**. The **features**, 'RM', 'LSTAT', and 'PTRATIO', give us quantitative information about each data point. The **target variable**, 'MEDV', will be the variable we seek to predict. These are stored in `features` and `prices`, respectively.

### 1.4.1 Implementation: Calculate Statistics

For your very first coding implementation, you will calculate descriptive statistics about the Boston housing prices. Since numpy has already been imported for you, use this library to perform the necessary calculations. These statistics will be extremely important later on to analyze various prediction results from the constructed model.

In the code cell below, you will need to implement the following: - Calculate the minimum, maximum, mean, median, and standard deviation of 'MEDV', which is stored in prices. - Store each calculation in their respective variable.

```
In [6]: # TODO: Minimum price of the data
        minimum_price = np.amin(prices)

        # TODO: Maximum price of the data
        maximum_price = np.amax(prices)

        # TODO: Mean price of the data
        mean_price = np.mean(prices)

        # TODO: Median price of the data
        median_price = np.median(prices)

        # TODO: Standard deviation of prices of the data
        std_price = np.std(prices)

        # Show the calculated statistics
        print("Statistics for Boston housing dataset:\n")
        print("Minimum price: ${}".format(minimum_price))
        print("Maximum price: ${}".format(maximum_price))
        print("Mean price: ${}".format(mean_price))
        print("Median price ${}".format(median_price))
        print("Standard deviation of prices: ${}".format(std_price))
```

Statistics for Boston housing dataset:

```
Minimum price: $105000.0
Maximum price: $1024800.0
Mean price: $454342.9447852761
Median price $438900.0
Standard deviation of prices: $165171.13154429474
```

### 1.4.2 Question 1 - Feature Observation

As a reminder, we are using three features from the Boston housing dataset: 'RM', 'LSTAT', and 'PTRATIO'. For each data point (neighborhood): - 'RM' is the average number of rooms among homes in the neighborhood. - 'LSTAT' is the percentage of homeowners in the neighborhood considered "lower class" (working poor). - 'PTRATIO' is the ratio of students to teachers in primary and secondary schools in the neighborhood.

**\*\* Using your intuition, for each of the three features above, do you think that an increase in the value of that feature would lead to an **increase** in the value of 'MEDV' or a **decrease** in the value of 'MEDV'? Justify your answer for each.\*\***

**Hint:** This problem can be phrased using examples like below.

\* Would you expect a home that has an 'RM' value (number of rooms) of 6 to be worth more or less than a home that has an 'RM' value of 7? \* Would you expect a neighborhood that has an 'LSTAT' value (percent of lower class workers) of 15 to have home prices be worth more or less than a neighborhood that has an 'LSTAT' value of 20? \* Would you expect a neighborhood that has an 'PTRATIO' value (ratio of students to teachers) of 10 to have home prices be worth more or less than a neighborhood that has an 'PTRATIO' value of 15?

**Answer:** 1. For a higher RM (number of rooms), I would expect to observe a higher MEDV (Price). The same can be seen in figure 1 when we plot Number of Rooms - Price Relation. 2. For a higher LSTAT (percent of lower class workers), I would expect to observe a lower MEDV (Price). The same can be seen in figure 2 when we plot % of lower class - Price Relation. 3. For a higher PTRATIO (ratio of student to teachers), I would expect to observe a lower MEDV (Price). The same can be seen in figure 2 when we plot Ratio of Student to Teacher - Price Relation.

Following are the plots justifying the answers:

```
In [7]: import matplotlib.pyplot as plt
```

```
# Number of Rooms - Price Relation
fig=plt.figure()
subfig=fig.add_subplot(1, 1, 1)
subfig.scatter(data['RM'], data['MEDV'])
plt.title("Prices - Number of rooms")
plt.xlabel("RM")
plt.ylabel("Prices")
plt.show()
print('1. The upward scatter plot clearly indicates that,
as the number of the rooms increase;
there is an increase in the price.\n\n')

# % of lower class - Price Relation
fig=plt.figure()
subfig=fig.add_subplot(1, 1, 1)
subfig.scatter(data['LSTAT'], data['MEDV'])
plt.title("Prices - % of low class Homeowners")
plt.xlabel("LSTAT")
plt.ylabel("Prices")
plt.show()
print('2. The downward scatter plot clearly indicates that,
as the poverty level increases;
there is a decrease in the price.\n\n')


# Ratio of Student to Teacher - Price Relation
fig=plt.figure()
subfig=fig.add_subplot(1, 1, 1)
```




```

subfig.scatter(data['PTRATIO'], data['MEDV'])
plt.title("Prices - Ratio of Students to Teachers")
plt.xlabel("PTRATIO")
plt.ylabel("Prices")
plt.show()
print('3. The downward scatter plot clearly indicates that,
as the ratio of student to teachers increase;
there is a decrease in the price.\n \n')


```



1. The upward scatter plot clearly indicates that,  
as the number of the rooms increase;  
there is a increase in the price.



2. The downward scatter plot clearly indicates that, as the poverty level increase; there is a decrease in the price.



3. The downward scatter plot clearly indicates that, as the ratio of student to teachers increase; there is a decrease in the price.

---

## 1.5 Developing a Model

In this second section of the project, you will develop the tools and techniques necessary for a model to make a prediction. Being able to make accurate evaluations of each model's performance through the use of these tools and techniques helps to greatly reinforce the confidence in your predictions.

### 1.5.1 Implementation: Define a Performance Metric

It is difficult to measure the quality of a given model without quantifying its performance over training and testing. This is typically done using some type of performance metric, whether it is through calculating some type of error, the goodness of fit, or some other useful measurement. For this project, you will be calculating the *coefficient of determination*,  $R^2$ , to quantify your model's performance. The coefficient of determination for a model is a useful statistic in regression analysis, as it often describes how "good" that model is at making predictions.

The values for R2 range from 0 to 1, which captures the percentage of squared correlation between the predicted and actual values of the **target variable**. A model with an R2 of 0 is no better than a model that always predicts the *mean* of the target variable, whereas a model with an R2 of 1 perfectly predicts the target variable. Any value between 0 and 1 indicates what percentage of the target variable, using this model, can be explained by the **features**. *A model can be given a negative R2 as well, which indicates that the model is **arbitrarily worse** than one that always predicts the mean of the target variable.*

For the performance\_metric function in the code cell below, you will need to implement the following: - Use r2\_score from sklearn.metrics to perform a performance calculation between y\_true and y\_predict. - Assign the performance score to the score variable.

```
In [8]: # TODO: Import 'r2_score'
        from sklearn.metrics import r2_score

        def performance_metric(y_true, y_predict):
            """ Calculates and returns the performance score between
                true and predicted values based on the metric chosen. """

            # TODO: Calculate the performance score between 'y_true' and 'y_predict'
            score = r2_score(y_true, y_predict)

            # Return the score
            return score
```

### 1.5.2 Question 2 - Goodness of Fit

Assume that a dataset contains five data points and a model made the following predictions for the target variable:

True Value	Prediction
3.0	2.5
-0.5	0.0
2.0	2.1
7.0	7.8
4.2	5.3

Run the code cell below to use the performance\_metric function and calculate this model's coefficient of determination.

```
In [9]: # Calculate the performance of this model
        score = performance_metric([3, -0.5, 2, 7, 4.2], [2.5, 0.0, 2.1, 7.8, 5.3])
        print("Model has a coefficient of determination, R^2, of {:.3f}.".format(score))
```

Model has a coefficient of determination, R<sup>2</sup>, of 0.923.

- Would you consider this model to have successfully captured the variation of the target variable?

- Why or why not?

**\*\* Hint: \*\*** The R2 score is the proportion of the variance in the dependent variable that is predictable from the independent variable. In other words: \* R2 score of 0 means that the dependent variable cannot be predicted from the independent variable. \* R2 score of 1 means the dependent variable can be predicted from the independent variable. \* R2 score between 0 and 1 indicates the extent to which the dependent variable is predictable. An \* R2 score of 0.40 means that 40 percent of the variance in Y is predictable from X.

**Answer:** The model has  $R^2$  value of 0.923(close to 1) which means that there is very high predictability indicating a strong correlation (93%). So I would agree that the model has successfully captured the variation of target variable. Anyhow, this deduction is made only using five data points. More data could further strengthen the conclusion.

### 1.5.3 Implementation: Shuffle and Split Data

Your next implementation requires that you take the Boston housing dataset and split the data into training and testing subsets. Typically, the data is also shuffled into a random order when creating the training and testing subsets to remove any bias in the ordering of the dataset.

For the code cell below, you will need to implement the following: - Use `train_test_split` from `sklearn.cross_validation` to shuffle and split the features and prices data into training and testing sets. - Split the data into 80% training and 20% testing. - Set the `random_state` for `train_test_split` to a value of your choice. This ensures results are consistent. - Assign the train and testing splits to `X_train`, `X_test`, `y_train`, and `y_test`.

```
In [10]: # TODO: Import 'train_test_split'
         from sklearn.cross_validation import train_test_split

         # TODO: Shuffle and split the data into training and testing subsets
         X_train, X_test, y_train, y_test = train_test_split(features, prices, test_size=0.2, ra

         # Success
         print("Training and testing split was successful.")
```

Training and testing split was successful.

### 1.5.4 Question 3 - Training and Testing

- What is the benefit to splitting a dataset into some ratio of training and testing subsets for a learning algorithm?

**Hint:** Think about how overfitting or underfitting is contingent upon how splits on data is done.

**Answer:** Splitting of dataset into training and testing subsets is required as we want to validate the predictive ability of the model on unseen data. While evaluating a predictive model on training data we don't know how well the model has generalized to new unseen data.

If we don't split our data, training accuracy is more which rewards overly complex models that overfit the training data. As we add more features, we can slowly reach a 100% accuracy but will not be able to generalize well.

Some problems with splitting the data into training and testing subsets: Testing accuracy can change a lot depending on which observation happens to be in the testing set. Hence we have to use techniques like k-fold cross-validation to solve this problem.

## 1.6 Analyzing Model Performance


In this third section of the project, you'll take a look at several models' learning and testing performances on various subsets of training data. Additionally, you'll investigate one particular algorithm with an increasing 'max\_depth' parameter on the full training set to observe how model complexity affects performance. Graphing your model's performance based on varying criteria can be beneficial in the analysis process, such as visualizing behavior that may not have been apparent from the results alone.

### 1.6.1 Learning Curves

The following code cell produces four graphs for a decision tree model with different maximum depths. Each graph visualizes the learning curves of the model for both training and testing as the size of the training set is increased. Note that the shaded region of a learning curve denotes the uncertainty of that curve (measured as the standard deviation). The model is scored on both the training and testing sets using R2, the coefficient of determination.

Run the code cell below and use these graphs to answer the following question.

```
In [11]: # Produce learning curves for varying training set sizes and maximum depths
vs.ModelLearning(features, prices)
```



### 1.6.2 Question 4 - Learning the Data

- Choose one of the graphs above and state the maximum depth for the model.
- What happens to the score of the training curve as more training points are added? What about the testing curve?
- Would having more training points benefit the model?

**Hint:** Are the learning curves converging to particular scores? Generally speaking, the more data you have, the better. But if your training and testing curves are converging with a score above your benchmark threshold, would this be necessary? Think about the pros and cons of adding more training points based on if the training and testing curves are converging.

**Answer:** Let's choose the second model where the maximum depth of the model is 3, this curve indicates an ideal scenario where the training curve (red line) and testing curve (green line) meet around 0.8; meaning it's not overly fitted (high variance) like the model with `max_depth = 10`, nor it's underfitted (high bias) like the model with `max_depth = 1`.

The graph has already flattened around 0.8, further adding training points would not have much impact on the curve. There won't be much impact on the testing curve as more training points are added.


More training points would help the model become ideal but it would work even without it, as the curves are well converged already.

### 1.6.3 Complexity Curves

The following code cell produces a graph for a decision tree model that has been trained and validated on the training data using different maximum depths. The graph produces two complexity curves — one for training and one for validation. Similar to the **learning curves**, the shaded regions of both the complexity curves denote the uncertainty in those curves, and the model is scored on both the training and validation sets using the `performance_metric` function.

**\*\* Run the code cell below and use this graph to answer the following two questions Q5 and Q6. \*\***

```
In [12]: vs.ModelComplexity(X_train, y_train)
```



#### 1.6.4 Question 5 - Bias-Variance Tradeoff

- When the model is trained with a maximum depth of 1, does the model suffer from high bias or from high variance?
- How about when the model is trained with a maximum depth of 10? What visual cues in the graph justify your conclusions?

**Hint:** High bias is a sign of underfitting(model is not complex enough to pick up the nuances in the data) and high variance is a sign of overfitting(model is by-hearting the data and cannot generalize well). Think about which model(depth 1 or 10) aligns with which part of the tradeoff.

**Answer:** When the model is trained with max\_depth is 1, the model suffers from high bias. In this case the model is generalising too much, this is a sign of underfitting, also the model is not complex enough.

When the max\_depth is 10, the model will suffer from high variance.

High bias models have a small or no gap between the training and validations score. In case of high variance models there is a substantial gap between the training and testing scores.

#### 1.6.5 Question 6 - Best-Guess Optimal Model

- Which maximum depth do you think results in a model that best generalizes to unseen data?
- What intuition lead you to this answer?



**\*\* Hint: \*\*** Look at the graph above Question 5 and see where the validation scores lie for the various depths that have been assigned to the model. Does it get better with increased depth? At what point do we get our best validation score without overcomplicating our model? And remember, Occams Razor states "Among competing hypotheses, the one with the fewest assumptions should be selected."

**Answer:** The maximum depth that would result in a model that best generalizes to unseen data would be 4.

There is not much gap between training and validation score. Therefore, it's not a case of overfitting. After this point the training and validation score diverges.

---

## 1.7 Evaluating Model Performance

In this final section of the project, you will construct a model and make a prediction on the client's feature set using an optimized model from `fit_model`.

### 1.7.1 Question 7 - Grid Search

- What is the grid search technique?
- How it can be applied to optimize a learning algorithm?

**\*\* Hint: \*\*** When explaining the Grid Search technique, be sure to touch upon why it is used, what the 'grid' entails and what the end goal of this method is. To solidify your answer, you can also give an example of a parameter in a model that can be optimized using this approach.

**Answer:** Grid Search is a technique to tune the hyper-parameters and find the best optimized model.

If we have a set of models to choose from (which differ from each other in their parameter values, which lie on a grid), we then train each of the models and evaluate it using cross-validation. We then select the model with best cross-validation score.

### 1.7.2 Question 8 - Cross-Validation

- What is the k-fold cross-validation training technique?
- What benefit does this technique provide for grid search when optimizing a model?

**Hint:** When explaining the k-fold cross validation technique, be sure to touch upon what 'k' is, how the dataset is split into different parts for training and testing and the number of times it is run based on the 'k' value.

When thinking about how k-fold cross validation helps grid search, think about the main drawbacks of grid search which are hinged upon **using a particular subset of data for training or testing** and how k-fold cv could help alleviate that. You can refer to the [docs](#) for your answer.

**Answer:**

k-fold cross-validation is a training technique where the training sample (a part of the original sample) is randomly partitioned into k equal sized sub-samples.

From the k sub-samples, a single sub-sample is retained as the validation data for testing the model, and the remaining k - 1 subsamples are used as training data. The cross-validation process is then repeated k times, with each of the k subsamples used exactly once as the validation data.

The k results can then be averaged to produce a single estimation as the overall result of the model called cross validation error. This Cross Validation error help us do the model selection. The selected model is again is validated with world data (the data which is part of original sample after the training sample). The result thus obtained gives us a better idea of how the model will work in a real world scenario.

The advantage of this method is that all observations are used for both training and validation, and each observation is used for validation exactly once and thus making the process more data-efficient reducing the waste too which is a major advantage in problems where the data samples is small.

### 1.7.3 Implementation: Fitting a Model

Your final implementation requires that you bring everything together and train a model using the **decision tree algorithm**. To ensure that you are producing an optimized model, you will train the model using the grid search technique to optimize the 'max\_depth' parameter for the decision tree. The 'max\_depth' parameter can be thought of as how many questions the decision tree algorithm is allowed to ask about the data before making a prediction. Decision trees are part of a class of algorithms called *supervised learning algorithms*.

In addition, you will find your implementation is using ShuffleSplit() for an alternative form of cross-validation (see the 'cv\_sets' variable). While it is not the K-Fold cross-validation technique you describe in **Question 8**, this type of cross-validation technique is just as useful!. The ShuffleSplit() implementation below will create 10 ('n\_splits') shuffled sets, and for each shuffle, 20% ('test\_size') of the data will be used as the *validation set*. While you're working on your implementation, think about the contrasts and similarities it has to the K-fold cross-validation technique.

Please note that ShuffleSplit has different parameters in scikit-learn versions 0.17 and 0.18. For the fit\_model function in the code cell below, you will need to implement the following: - Use `DecisionTreeRegressor` from `sklearn.tree` to create a decision tree regressor object. - Assign this object to the 'regressor' variable. - Create a dictionary for 'max\_depth' with the values from 1 to 10, and assign this to the 'params' variable. - Use `make_scorer` from `sklearn.metrics` to create a scoring function object. - Pass the performance\_metric function as a parameter to the object. - Assign this scoring function to the 'scoring\_fnc' variable. - Use `GridSearchCV` from `sklearn.grid_search` to create a grid search object. - Pass the variables 'regressor', 'params', 'scoring\_fnc', and 'cv\_sets' as parameters to the object. - Assign the GridSearchCV object to the 'grid' variable.

```
In [13]: # TODO: Import 'make_scorer', 'DecisionTreeRegressor', and 'GridSearchCV'
         from sklearn.metrics import make_scorer
         from sklearn.tree import DecisionTreeRegressor
         from sklearn.grid_search import GridSearchCV

         def fit_model(X, y):
             """ Performs grid search over the 'max_depth' parameter for a
                 decision tree regressor trained on the input data [X, y]. """

             # Create cross-validation sets from the training data
             # sklearn version 0.18: ShuffleSplit(n_splits=10, test_size=0.1, train_size=None, r
             # sklearn version 0.17: ShuffleSplit(n, n_iter=10, test_size=0.1, train_size=None,
```

```

cv_sets = ShuffleSplit(X.shape[0], n_iter = 10, test_size = 0.20, random_state = 0)

# TODO: Create a decision tree regressor object
regressor = DecisionTreeRegressor()

# TODO: Create a dictionary for the parameter 'max_depth' with a range from 1 to 10
params = {'max_depth': list(range(1,11))}

# TODO: Transform 'performance_metric' into a scoring function using 'make_scorer'
scoring_fnc = make_scorer(performance_metric)

# TODO: Create the grid search cv object --> GridSearchCV()
# Make sure to include the right parameters in the object:
# (estimator, param_grid, scoring, cv) which have values 'regressor', 'params', 'scoring_fnc', 'cv_sets'
grid = GridSearchCV(regressor, params, scoring = scoring_fnc, cv = cv_sets)

# Fit the grid search object to the data to compute the optimal model
grid = grid.fit(X, y)

# Return the optimal model after fitting the data
return grid.best_estimator_

```

### 1.7.4 Making Predictions

Once a model has been trained on a given set of data, it can now be used to make predictions on new sets of input data. In the case of a *decision tree regressor*, the model has learned *what the best questions to ask about the input data are*, and can respond with a prediction for the **target variable**. You can use these predictions to gain information about data where the value of the target variable is unknown — such as data the model was not trained on.

### 1.7.5 Question 9 - Optimal Model

- What maximum depth does the optimal model have? How does this result compare to your guess in **Question 6**?

Run the code block below to fit the decision tree regressor to the training data and produce an optimal model.

```

In [14]: # Fit the training data to the model using grid search
         reg = fit_model(X_train, y_train)

         # Produce the value for 'max_depth'
         print("Parameter 'max_depth' is {} for the optimal model.".format(reg.get_params()['max_depth']))

```

Parameter 'max\_depth' is 4 for the optimal model.

**\*\* Hint: \*\*** The answer comes from the output of the code snipped above.

**Answer:** The optimal depth for the model is 4.

### 1.7.6 Question 10 - Predicting Selling Prices

Imagine that you were a real estate agent in the Boston area looking to use this model to help price homes owned by your clients that they wish to sell. You have collected the following information from three of your clients:

Feature	Client 1	Client 2	Client 3
Total number of rooms in home	5 rooms	4 rooms	8 rooms
Neighborhood poverty level (as %)	17%	32%	3%
Student-teacher ratio of nearby schools	15-to-1	22-to-1	12-to-1

- What price would you recommend each client sell his/her home at?
- Do these prices seem reasonable given the values for the respective features?

**Hint:** Use the statistics you calculated in the **Data Exploration** section to help justify your response. Of the three clients, client 3 has the biggest house, in the best public school neighborhood with the lowest poverty level; while client 2 has the smallest house, in a neighborhood with a relatively high poverty rate and not the best public schools.

Run the code block below to have your optimized model make predictions for each client's home.

```
In [15]: # Produce a matrix for client data
        client_data = [[5, 17, 15], # Client 1
                        [4, 32, 22], # Client 2
                        [8, 3, 12]]  # Client 3

        # Show predictions
        for i, price in enumerate(reg.predict(client_data)):
            print("Predicted selling price for Client {}'s home: ${:,.2f}".format(i+1, price))
```

Predicted selling price for Client 1's home: \$391,183.33

Predicted selling price for Client 2's home: \$189,123.53

Predicted selling price for Client 3's home: \$942,666.67

**Answer:** Minimum price: \$105,000.00

Maximum price: \$1,024,800.00

Mean price: \$454,342.94

Median price \$438,900.00

Standard deviation of prices: \$165,340.28

**Recommended Prices:**

Client 1's home: \$391,183.33 (Seems reasonable due to the average poverty level and average student-teacher ratio)

Client 2's home: \$189,123.53 (Seems reasonable due to the high poverty level and high student-teacher ratio.)

Client 3's home: \$942,666.67 (Seems reasonable as the number of rooms are more also the poverty level and student-teacher ratio is low)

### 1.7.7 Sensitivity

An optimal model is not necessarily a robust model. Sometimes, a model is either too complex or too simple to sufficiently generalize to new data. Sometimes, a model could use a learning algorithm that is not appropriate for the structure of the data given. Other times, the data itself could be too noisy or contain too few samples to allow a model to adequately capture the target variable — i.e., the model is underfitted.

**Run the code cell below to run the `fit_model` function ten times with different training and testing sets to see how the prediction for a specific client changes with respect to the data it's trained on.**

```
In [16]: vs.PredictTrials(features, prices, fit_model, client_data)
```

```
Trial 1: $391,183.33
Trial 2: $419,700.00
Trial 3: $415,800.00
Trial 4: $420,622.22
Trial 5: $418,377.27
Trial 6: $411,931.58
Trial 7: $399,663.16
Trial 8: $407,232.00
Trial 9: $351,577.61
Trial 10: $413,700.00
```

```
Range in prices: $69,044.61
```

### 1.7.8 Question 11 - Applicability

- In a few sentences, discuss whether the constructed model should or should not be used in a real-world setting.

**Hint:** Take a look at the range in prices as calculated in the code snippet above. Some questions to answering: - How relevant today is data that was collected from 1978? How important is inflation? - Are the features present in the data sufficient to describe a home? Do you think factors like quality of appliances in the home, square feet of the plot area, presence of pool or not etc should factor in? - Is the model robust enough to make consistent predictions? - Would data collected in an urban city like Boston be applicable in a rural city? - Is it fair to judge the price of an individual home based on the characteristics of the entire neighborhood?

**Answer:** The model should not be used in the real-world setting because: - The data used to train the model is quite old and thus might not be relevant today. - The selected domain(real state housing) is geographic dependent problem, but there is no feature about it. - The model also appears to be too sensitive, as it show wide variability in pricing. - The data used is from Boston, which is an urban setup, this might not be case for the area we are trying to predict. - Although neighborhood plays an important role, but there may be cases where the house is totally independent, the judgement in that case may be totally off.

**Note:** Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython

Notebook as an HTML document. You can do this by using the menu above and navigating to

**File -> Download as -> HTML (.html).** Include the finished document along with this notebook as your submission.